

# Data Analysis 3

Decision trees, Rule based classification, Probabilistic classification

---

Jan Platoš

November 2, 2020

Department of Computer Science  
Faculty of Electrical Engineering and Computer Science  
VŠB - Technical University of Ostrava

# Table of Contents

Decision trees

Rule-based classification

Naïve Bayes Classifier

Classification Assesment

# Decision trees

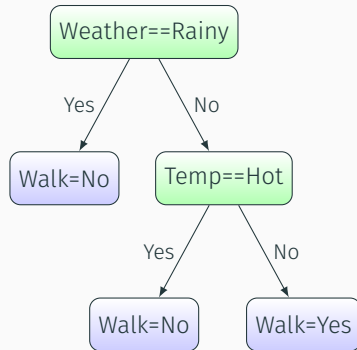
---

# Decision Trees

- Classification is modeled using hierarchical decisions on the features that are arranged in tree-like structure.
- The decision at a particular node, called split criterion, is a relational condition on one or more features and their values.
- The goal is to identify a split criterion that minimizes the mixing of classes in each branch.
- Works on binary, numeric and categorical attributes.
- Each sub-space (region) is recursively split until terminal conditions are reached.
- Univariate or Multivariate split is possible.

# Decision Trees

Weather	Temp	Walk?
Sunny	Cold	Yes
Sunny	Warm	Yes
Sunny	Hot	No
Cloudy	Cold	Yes
Cloudy	Warm	Yes
Cloudy	Hot	No
Rainy	Cold	No
Rainy	Warm	No
Rainy	Hot	No



## Split Criteria:

- The goal is to maximize separation of the different classes among the children nodes.
- Binary attribute – only one type of split is possible.
- Categorical attribute with  $r$  values
  - $r$ -way split,
  - binary split on  $2^r - 1$  possibilities (all combinations except  $\emptyset$ ),
  - binary split on  $r$  possibilities (one-to-rest strategy).
- Numeric attribute
  - A split is made between two values with  $<$  or  $\leq$  relation.
  - All values or selected values only may be tested.

Definitions:

- $S$  is a set of points in a branch of a tree.
- $|S|$  is size of the set (number of points in a set).
- $r$ -way split has  $r$  subsets  $S_1, \dots, S_r$  of set  $S$ .
- $k$  is the number of classes.

Error rate:

- On a set:

$$Err(S) = 1 - p$$

- where the  $p$  is a fraction of points that belongs to the dominant class from  $S$ .
- On  $r$ -way split:

$$Err(S \Rightarrow S_1, \dots, S_r) = \sum_{i=1}^r \frac{|S_i|}{|S|} (1 - p)$$



Gini index:

- On a set:

$$G(S) = 1 - \sum_{j=1}^k p_j^2$$

- where the  $p_j$  is a fraction of points that belongs to the class  $j$  from  $S$ .
- On  $r$ -way split:

$$G(S \Rightarrow S_1, \dots, S_r) = \sum_{i=1}^r \frac{|S_i|}{|S|} G(S_i)$$

# Decision Trees - Split Criteria

Entropy:

- On a set:

$$E(S) = - \sum_{j=1}^k p_j \log_2(p_j)$$

- where the  $p_j$  is a fraction of points that belongs to the class  $j$  from  $S$ .
- On  $r$ -way split:

$$E(S \Rightarrow S_1, \dots, S_r) = \sum_{i=1}^r \frac{|S_i|}{|S|} E(S_i)$$

## Decision Trees - Split Criteria - IRIS Example

## Decision Trees - Split Criteria - IRIS Example

## Decision Trees - Split Criteria - IRIS Example

## Decision Trees - Split Criteria - IRIS Example

## Decision Trees - Split Criteria - IRIS Example

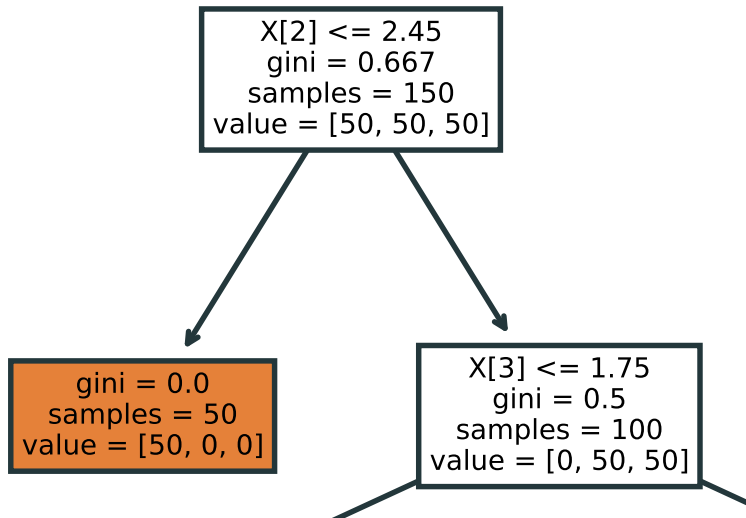
- index = 0, Sepal Length  $< 5.45$ , Gini = 0.44
- index = 1, Sepal Width  $< 3.35$ , Gini = 0.54
- index = 2, Petal Length  $< 2.45$ , Gini = 0.33
- index = 3, Petal Width  $< 0.80$ , Gini = 0.33

## Decision Trees - Split Criteria - IRIS Example

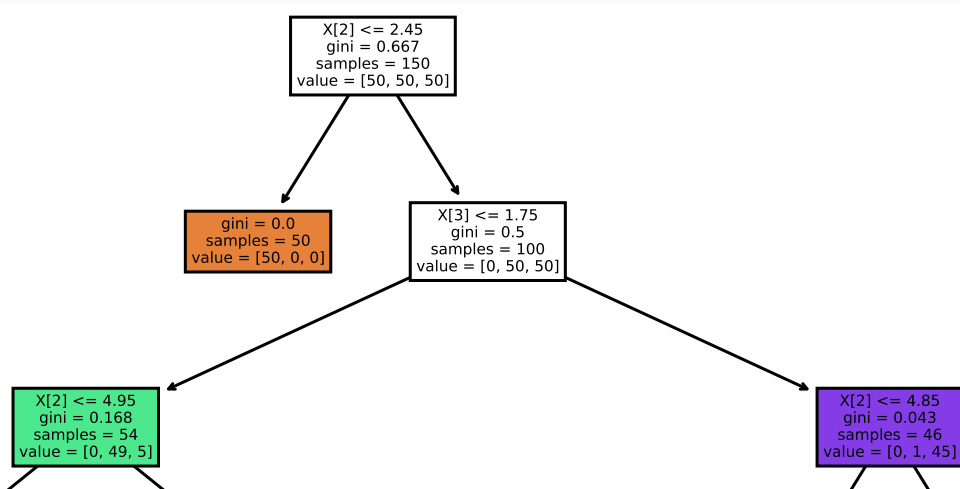
- index = 0, Sepal Length  $< 5.45$ , Gini = 0.44
- index = 1, Sepal Width  $< 3.35$ , Gini = 0.54
- index = 2, Petal Length  $< 2.45$ , Gini = 0.33
- index = 3, Petal Width  $< 0.80$ , Gini = 0.33



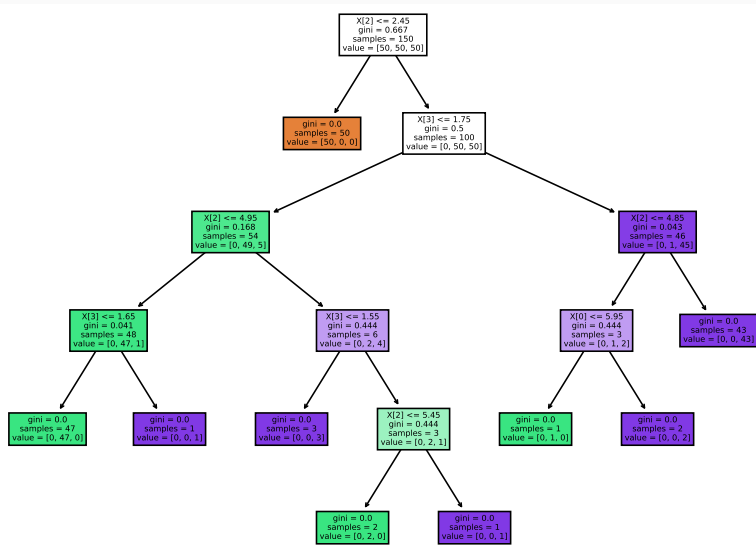
## Decision Trees - Split Criteria - IRIS Example



# Decision Trees - Split Criteria - IRIS Example



# Decision Trees - Split Criteria - IRIS Example



Stopping criterion:

- Very difficult to stop during the tree growth.
- Single class in a leaf node is the final condition.
- Such tree has 100% precision on Training data.
- But, such tree is over-fitted (unable to generalize to unseen data).
- Over-fitting is done by lower nodes with less number of points.

## Pruning:

- Shallow trees are more preferable if they produce the same error on training data.
- Nodes/Trees are evaluated using a criterion that penalizes the more complex tree without satisfactory improvement in precision.
- Usually a holdout set (e.g. 20% of training set) is used for pruning.
- A node is pruned if its removal improves the precision on the holdout.
- A leaf node is pruned iteratively until no node should be removed.

## Rule-based classification

---

- A generalization of the Decision Trees.
- A set of rules in a form:

*IF Condition THEN Conclusion*

- *Condition* or *Antecedent* is a combination of relational, set and logical operators over features.
- *Conclusion* or *Consequent* is a class label.
- A rule cover the training instance is the condition match the instance.

Rule types:

- Mutually exclusive rules
  - Each rule covers disjoint set of instances.
  - Each instance trigger at most one rule.
- Exhaustive rules
  - The entire data space is covered by at least one rule.
  - Simple exhaustive rule assign dominant class do anything (catch-all).
- Non mutually exclusive rules brings problems with rule evaluation.



Rule ordering:

- Ordered rules
  - Rules are ordered by priority, such as quality measure.
  - Rules may be ordered by class-based principle.
  - Only the first triggered rule vote, its consequent is the result.
  - The rare classes are usually ordered first.
- Unordered rules
  - There is no priority on rules.
  - The dominant class of the all triggered rules is selected.
  - Simplifies the learning phase.

Rule generation:

- The goal is to generate rules that covers the instances from the training data.
- Two major algorithm exists:
  - Generation using Decision Trees.
  - Sequential Covering Algorithm.

# Rule-based classification - Rule generation

Rule generation using Decision Trees:

- Trees are used for generation of the rules.
- Each leaf node represent one rule with its sequence of splits that lead to this leaf from root.
- The pruning is not made on tree, but on rules.
- Each rule is processed separately and pruned to get the most precise rule on the holdout set.
- The pruning process is more flexible because any part of the antecedent may be pruned.
- Duplicate rules are removed.
- The rules after pruning are not mutually exclusive.
- The ordering of the rules is necessary.
- Rare classes and less complex rules or rules with less false positives are prioritized.

## Sequential Covering Algorithm:

- An algorithm for creation of ordered set of rules.
- An 2-step iterative algorithm:
  - **Learn-one-rule** – select particular class and determine the “best” rule from the current training instances  $S$  with this class as a consequent. Add this rule to the bottom of the ordered rule list.
  - **Prune training data** – Remove training instances in  $S$  that are covered by the rule generated in previous step. The detection is based on the antecedent only, that consequent of the instances is ignored.

The ordering of the generated rules:

- Class-based ordering
  - All rules for particular class are put together.
  - Rare classes may be prioritized.
  - All rules for this particular class are generated continuously, until a termination criterion is met.
  - For  $k$ -class problem,  $k - 1$  rule sets are generated and the final catch-all rule covers the last class.
- Quality-based ordering
  - The rules are selected according to a measure, such as confidence or support.
  - The catch-all rule corresponds to the dominant class among remaining instances.
  - The quality is very difficult to measure.

Learn-one-rule step:

- Iterative algorithm that grows a rule with best conjunct according the quality measure.
- The simplest quality is the precision/accuracy.
- Each split choice (conjunct) is evaluated the same was as it is in trees.
- Several best options may be maintained to reduce the possibility of the mistakes and suboptimal rules.
- The ideal quality measure must combine accuracy and coverage, e.g. Laplace smoothing, like-hood ratio statistics, FOIL information gain.

Rule pruning:

- An Minimum description length (MDL) principle is one option.
- A penalty based on MDL may be used in rule-growth phase.
- An holdout set is another good principle.
- A greedy algorithm may be used for conjunct evaluation.

# Naïve Bayes Classifier

---



- Based on the Bayes theorem for conditional probabilities.

$$P(D|E) = \frac{P(E|D)P(D)}{P(E)}$$

- This theorem is useful when it is hard to estimate  $P(D|E)$  but others probabilities are easy to get from input data.
- When  $E$  is a single attribute, everything is simple.
- When  $E$  is complex,  $P(E|D)$  may be missing in the data or appear only few times.

- Let  $C$  be a class variable.
- Let  $\bar{X}$  is a  $d$ -dimensional instance  $\bar{X} = (a_1, \dots, a_d)$ .
- Let the random  $d$ -dimensional variable is  $\bar{X} = (x_1, \dots, x_d)$ .
- The goal is to estimate  $P(C = c | \bar{X} = (a_1, \dots, a_d))$  or  $P(C = c | x_1 = a_1, \dots, x_d = a_d)$  resp.

$$P(C = c | x_1 = a_1, \dots, x_d = a_d) = \frac{P(x_1 = a_1, \dots, x_d = a_d | C = c)P(C = c)}{P(x_1 = a_1, \dots, x_d = a_d)}$$

- The denominator is independent of the class and may be removed.

# Naïve Bayes Classifier

- The estimation of  $P(x_1 = a_1, \dots, x_d = a_d | C = c)$  is crucial and difficult.
- The Naïve approach assumes that the features are independent!!!
- Then

$$P(x_1 = a_1, \dots, x_d = a_d | C = c) = \prod_{j=1}^d P(x_j = a_j | C = c)$$

$$P(x_j = a_j | C = c) = \frac{q(a_j, c) + \alpha}{r(c) + \alpha \cdot m_j}$$

Where

- $q(a_j, c)$  is a fraction of records with class  $c$  and
- $r(c)$  is a fraction of records with class  $c$
- $\alpha$  is a small value
- $m_j$  is number of distinct values of  $j$ -th attribute.

- Finally:

$$P(C = c | x_1 = a_1, \dots, x_d = a_d) \approx P(C = c) \prod_{j=1}^d P(x_j = a_j | C = c)$$

- or

$$P(C = c | x_1 = a_1, \dots, x_d = a_d) \approx P(C = c) \prod_{j=1}^d \frac{q(a_j, c) + \alpha}{r(c) + \alpha \cdot m_j}$$

# Naïve Bayes Classifier - Example

Name	Age	Salary	Donor?
Nancy	21	37,000	N
Jim	27	41,000	N
Allen	43	61,000	Y
Jane	38	55,000	N
Steve	44	30,000	N
Peter	51	56,000	Y
Sayani	53	70,000	Y
Lata	56	74,000	Y
Mary	59	25,000	N
Victor	61	68,000	Y
Dale	63	51,000	Y

- Assume a rule: Age>50 AND Salary>50
- $P(\text{Donor} = \text{Yes}) = 6/11$ 
  - $P(\text{Age} > 50 | \text{Donor} = \text{Yes}) = 5/6$
  - $P(\text{Salary} > 50,000 | \text{Donor} = \text{Yes}) = 6/6$
- $P(\text{Donor} = \text{No}) = 5/11$ 
  - $P(\text{Age} > 50 | \text{Donor} = \text{No}) = 1/5$
  - $P(\text{Salary} > 50,000 | \text{Donor} = \text{No}) = 1/5$
- $P(\text{Donor} = \text{YES}) = 6/11 * 5/6 * 1 = 6/11 * 5/6 = 5/11$
- $P(\text{Donor} = \text{NO}) = 5/11 * 1/5 * 1/5 = 5/11 * 1/25 = 1/55$

# Naïve Bayes Classifier - Summary

- Numeric values
  - Discretization is possible but it may affect the precision.
  - Direct data-driven estimation of the probability distribution is more suitable.
  - A proper distribution has to be selected, usually Gaussian is taken.
  - A mean and variance is extracted from the data.
- The naïve assumption
  - The independence is usually not true in real data.
  - The more complex estimation of the probability are not precise when dimension increases.
  - The naïve approach is precise enough.

# Classification Assessment

---

- How do we quantify the accuracy of the given classification model?
- These methods has several applications: evaluation of the classification effectiveness, comparing different models, selecting the best model for a particular data set, parameter tuning and advanced meta-algorithms (ensemble).
- The issues related to this task may be divided into two categories:
  1. Methodological issues
    - The proper division of the labeled dataset into training and testing part.
    - This choice has direct impact on the evaluation process (overestimation or underestimation).
    - Several approaches are possible (holdout, bootstrap, cross-validation).
  2. Quantification issues
    - These methods are associated with the providing numerical measure for the quality with respect of the methodological issues.
    - Several methods output direct measure.
    - Other methods quantify the relative performance of classifiers.



## Methodological Issues

- These methods defines the partitioning of the ground-truth data for classification evaluation.
- The using of the same data for training and testing is not possible due to over-fitting and overestimation.
- In practice, the input data should be divided into three parts:
  - the model-building part of the labeled data
  - validation part of the labeled data
  - testing data.
- The validation part is used for parameter tuning or model selection.
- When the parameter tuning is done, the model is reconstructed on the whole dataset.
- The knowledge from the testing dataset should not be used in parameter tuning.

## Holdout

- The labeled data is randomly divided into two disjoint sets (training and testing).
- Typically 60% to 75% is used for training set.
- This partition may be repeated several times to get the final estimation.
- The over-presented samples in the training set are under-presented in the testing sets.
- Due to not using of the whole data set for training the estimation are pessimistic.
- By repeating the process over  $b$  different holdout samples the mean and the variance of the error estimates may be determined.
- These information may be used for building the confidence intervals on the error.
- In case of imbalanced data, an independent sampling (for each class separately) have to be used to ensure the similarity between whole dataset and the testing dataset.

## Cross-Validation

- The data is divided into  $m$  disjoint subsets of equal size  $n/m$ .
- A typical choice for  $m$  is around 10.
- One segment is used as a testing set the the remaining  $m - 1$  as a training set.
- This process is repeated by selection each of the  $m$  subsets as a testing sets.
- The average accuracy over the  $m$  different test sets is reported.
- The size of the training set is  $(m - 1) * n/m$ .
- When  $m$  is chosen large, the training set size is close to the whole dataset and the reported prediction is very close to the whole data set.
- The estimate of the accuracy tends to be highly representative but pessimistic.
- A special case is when  $m = n$ , this is called a *leave-one-out* cross-validation.
- *Stratified cross-validation* uses proportional representation of each class in the different folds and usually provides less pessimistic results.

## Bootstrap

- The labeled data are sampled uniformly with replacement to create a training set that may contain a duplicates.
- The labeled data of size  $n$  is sampled  $n$  times with replacement.
- The probability that a particular data point is not included in a sample is given by  $(1 - 1/n)$
- The probability that the point is not included in  $n$  samples is then  $(1 - 1/n)^n$ .
- For large values of  $n$  the expression is approximately  $1/e$ .
- The fraction of labeled points included included at least once in the dataset is  $1 - 1/e = 0.632$ .
- The training model is constructed on the bootstrapped sample with duplicates.
- The overall accuracy is computed using the whole dataset.
- The estimate is highly optimistic due to large overlap between training and testing set.

## Quantification Issues

- When the output of the classifier is in the form of a class label the prediction value is compared to the ground-truth.
- When the output of the classifier is in the form of a numerical score for each labeling possibility the label with highest score imply greater likelihood to a particular class.

## Output as Class Labels

- *Accuracy* - the fraction of test instances in which the predicted value matched the ground-truth value.
- *Cost-sensitive accuracy*
  - Not all cases are equally important in all scenarios while comparing the accuracy, e.g. Imbalanced data, ill vs. healthy patients, etc.
  - This is frequently quantified by imposing different costs  $c_1, \dots, c_k$  on the misclassification on the different classes.
  - Let  $n_1, \dots, n_k$  be the number of test instances belonging to each class.
  - Let  $a_1, \dots, a_k$  be the accuracies (expressed as a fraction) on the subset of test instances belonging to each class.
  - The overall accuracy  $A$  can be computed as a weighted combination of the accuracies over the individual labels:

$$A = \frac{\sum_{i=1}^k c_i n_i a_i}{\sum_{i=1}^k c_i n_i}$$

## Output as Numerical Score

- For simplicity, we will consider the two class classification problem.
- The numerical score provides more flexibility in evaluating the overall trade-off between labeling a varying number of data points as positives.
- The different setting of the threshold leads to different models.
- When the threshold is set too aggressive, the algorithm will miss the true-positives and false negatives.
- When threshold is too relaxed the algorithm produces many false-positives (false negatives).
- The correct threshold is not known a priori, but depends on the data.

- For any given threshold  $t$  on the predicted positive-class score the declared positive class set is denoted by  $S(t)$ .
- The size of the  $S(t)$  changes with the changes of the  $t$ .
- The  $G$  represents the true set (ground-truth) of positive instances.
- The *Precision* is defined as the percentage of reported positives that truly turn out to be positive

$$Precision(t) = 100 * \frac{|S(t) \cap G|}{|S(t)|}$$

- The value of *Precision*( $t$ ) is not necessarily monotonic in  $t$  because both numerator and denominator may change with  $t$  differently.



## Classification Assessment - Quantification Issues - Output as Numerical Score

- The *recall* is correspondingly defined as the percentage of ground-truth positives that have been reported as positive at threshold  $t$ .

$$Recall(t) = 100 * \frac{|S(t) \cap G|}{|G|}$$

- The natural trade-off between precision and recall exists, but it is not necessarily monotonic.
- The  $F_1$ -measure summarizes both precision and the recall.

$$F_1(t) = \frac{2 \cdot Precision(t) \cdot Recall(t)}{Precision(t) + Recall(t)}$$

- The  $F_1$ -measure provides better quantification than precision or recall, but it still depends on the  $t$ .
- The entire trade-off between recall and precision may be investigated by plotting these values with respect to the threshold  $t$ .

## ROC curve

- ROC curve is a different method for evaluating the trade-off which is more intuitive.
- The *true-positive rate* is the same as the recall:

$$TPR(t) = Recall(t) = 100 * \frac{|S(t) \cap G|}{|G|}$$

- The false positive rate is the percentage of the falsely reported positives out of the ground-truth negatives.

$$FPR(t) = 100 * \frac{|S(t) - G|}{|D - G|}$$

- The ROC curve is defined by plotting the  $FPR(t)$  on the x-axis and  $TPR(t)$  on the y-axis for varying values of  $t$ .

ROC curve cont.

- The ROC curve has always points  $(0, 0)$  and  $(100, 100)$ .
- The random classifier is expected to exhibit performance along the diagonal.
- The *lift* above the diagonal provides the idea about accuracy of the approach.
- The area below the ROC curve provides a concrete quantitative evaluation of the effectiveness of the particular method.

Questions?